

R Functions

Andy Hsu

All About Functions in R

Functions are the way we get stuff done in R. We call a function to read data, compute, plot, and do just about anything in R.

R makes writing our own function accessible, but it is important to understand the fundamentals and write a functioning snippet of code before diving into function creation.

Starting With a Snippet

To start, we will grade a class of student assignments. The first exercise will be with a small sample of 3 students.

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

To properly calculate the grade, we want to drop the students' lowest scores and calculate the average score afterward. We can do that with the following code.

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

To deal with Na values, we can use the argument `trim=` to check for Na values and set them to 0.

```
noNaTemp <- student2
noNaTemp[is.na(noNaTemp)] <- 0
```

Putting it all together, the final snippet of code is as follows. Note that the filtering of Na values needs to occur before dropping the lowest score.

```
noNaTemp <- student3
noNaTemp[is.na(noNaTemp)] <- 0
mean(noNaTemp[-which.min(noNaTemp)])
```

```
[1] 12.85714
```

Creating a Function

To package this into a function with its proper arguments, we can use `function(x)`, where `x` is the argument.

```
grade <- function(x) {
  # Set NA values to 0
  x[is.na(x)] <- 0
  # Drop the lowest score and take the mean
  mean(x[-which.min(x)])
}
```

Now, when we call the function with the argument of a scoreset that we want graded, R returns us the proper grade.

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

Using the `apply()` Function

Now, we want to use this function to grade multiple students at once. We first obtain the data frame of scores.

```
gradebook <- read.csv("https://tinyurl.com/gradeinput",row.names=1)
gradebook
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	NA	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	NA
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	NA	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

To allow the function to read multiple students at once from a data frame, we use the `apply()` function. The arguments for this function are: `X`, which specifies a dataset, `MARGIN`, which specifies how the function should be applied (i.e. by row, by column), and `FUN`, which specifies which function to be applied.

```
apply(gradebook,1,grade)
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

To find the highest scoring student in the class, we can take the `which.max()` of the result.

```
which.max(apply(gradebook,1,grade))
```

```
student-18  
18
```

Looking at averages across each assignment, we can see that the hardest assignment appears to be HW2.

```
y <- gradebook  
y[is.na(y)] <- 0  
apply(y,2,mean)
```

```
hw1 hw2 hw3 hw4 hw5  
89.00 72.80 80.80 85.15 79.25
```

To find which assignment was best correlated with score, we can call the `cor()` function within an `apply()`.

```
grades <- apply(gradebook,1,grade)  
y <- gradebook  
y[is.na(y)] <- 0  
apply(y,2,cor,grades)
```

```
hw1 hw2 hw3 hw4 hw5  
0.4250204 0.1767780 0.3042561 0.3810884 0.6325982
```

Reading the highest correlation coefficient, we find that HW5 was the most indicative of a student's score.

And those were some of the basics of writing your own functions in R!