

Breast Cancer Mini Project

Andy Hsu

Analysis of a Cancer Dataset

In this project, we will analyze a dataset of breast cancer biopsies from fine needle aspiration (FNA).

Getting Started

We will start this project by downloading the dataset, setting patient IDs to the row names.

```
wisc.df <- read.csv("wisconsin_cancer.csv", row.names=1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585

84348301	0.2597		0.09744	0.4956	1.1560	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
84358402	94.44	0.011490	0.02461	0.05688		0.01885
843786	27.19	0.007510	0.03345	0.03672		0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
84358402	0.01756		0.005115	22.54		16.67
843786	0.02165		0.005082	15.47		23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0	0.1622			0.6656
842517	158.80	1956.0	0.1238			0.1866
84300903	152.50	1709.0	0.1444			0.4245
84348301	98.87	567.7	0.2098			0.8663
84358402	152.20	1575.0	0.1374			0.2050
843786	103.40	741.6	0.1791			0.5249
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654			0.4601
842517	0.2416		0.1860			0.2750
84300903	0.4504		0.2430			0.3613
84348301	0.6869		0.2575			0.6638
84358402	0.4000		0.1625			0.2364
843786	0.5355		0.1741			0.3985
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

We will also separate the diagnosis, as it will not be used until later.

```
wisc.data <- wisc.df[,-1]
diagnosis <- wisc.df[,1]
```

Running commands to gain some basic information on the dataset, we find that there are 569 observations, 212 malignant diagnoses, and 10 variables suffixed with “_mean”. Importantly, we can use the `grep()` function to match the `colnames()` to the pattern of “_mean”, then enter the function `length()` to count the total.

```
dim(wisc.data)
```

```
[1] 569 30
```

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

```
length(grep("_mean", (colnames(wisc.data))))
```

```
[1] 10
```

Performing PCA

Before we can begin with PCA, we need to check that PCA can be applied to the dataset.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03

compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Now, we can run a PCA using the `prcomp()` function.

```
wisc.pr <- prcomp(wisc.data,scale=T)
summary(wisc.pr)
```

Importance of components:

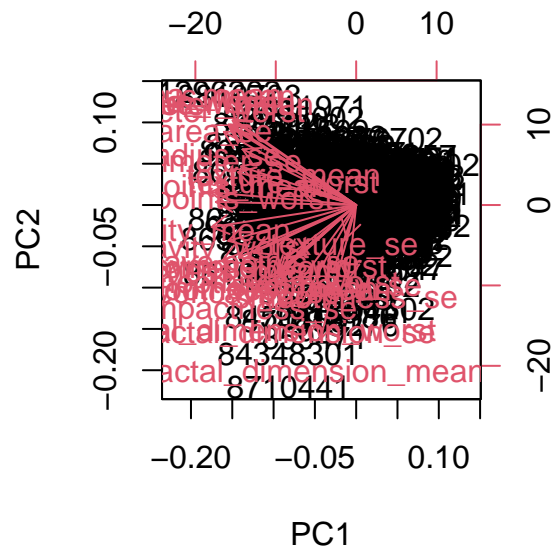
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Reading off the summary table, we can see that PC1 accounts for 44.3% of the original variance. Additionally, we need the first 3 PCs to reach 70% variance and 7 PCs to reach 90% variance.

Interpreting PCA

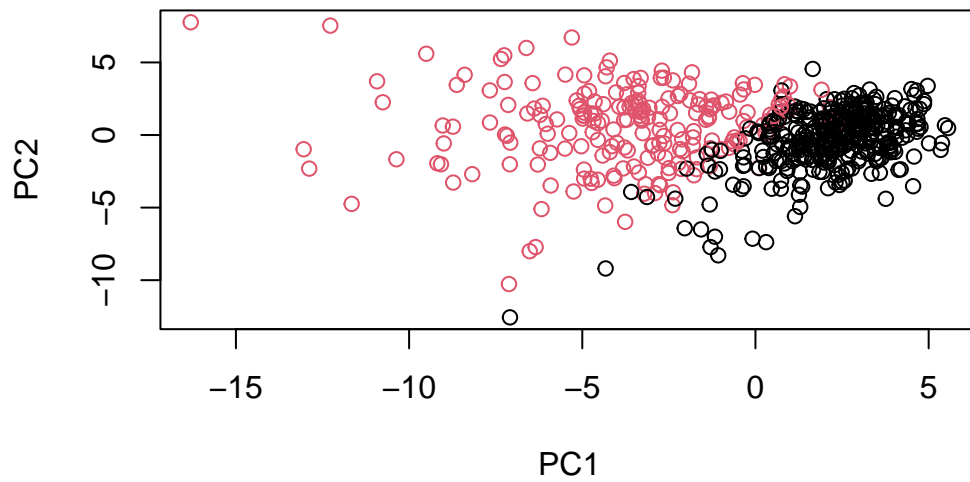
To take a look at our dataset, we can try to plot a biplot.

```
biplot(wisc.pr)
```



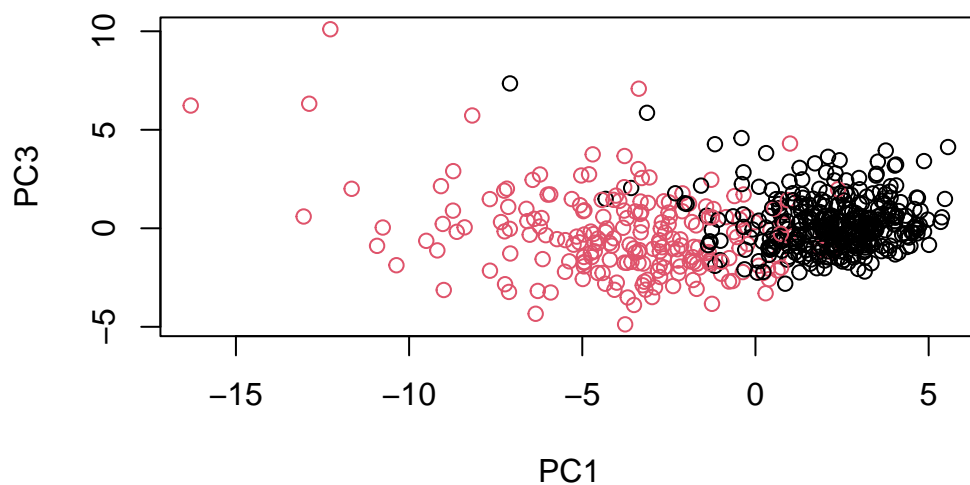
This plot is impossible to read, though. The overlapping text makes it too hard to tell where any datapoint is. Perhaps this plot would be better off with points labeled as dots rather than text.

```
plot(wisc.pr$x, col = (diagnosis=="M")+1, xlab = "PC1", ylab = "PC2")
```



To view this same plot with PC1 against PC3, we can modify the code. Notice that the distribution of the first plot shows a clearer separation between the 2 groups.

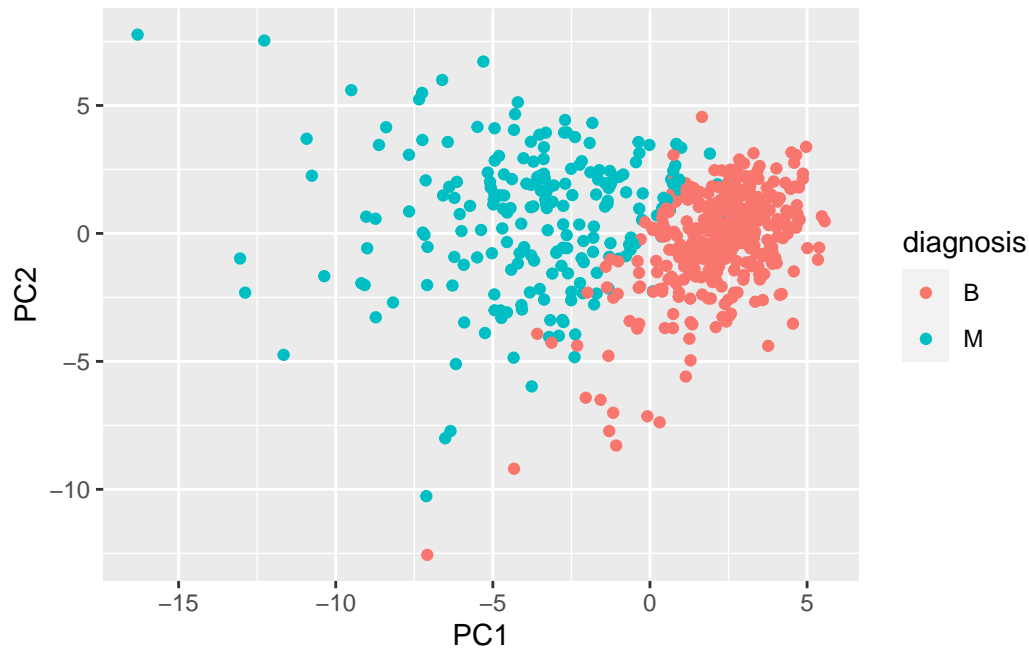
```
plot(wisc.pr$x[,c(1,3)], col = (diagnosis=="M")+1, xlab = "PC1", ylab = "PC3")
```



It's now time to plot our results using ggplot2. First, we need to convert our data to a dataframe, including the diagnosis, then we can plot.

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

library(ggplot2)
ggplot(df, aes(PC1,PC2,col=diagnosis)) +
  geom_point()
```

Plotting Variance on a Scree Plot


We can also visualize our variance against the number of PCs via a scree plot. First, we calculate the variance of each PC.

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

Next, we can calculate the proportion of the total variance explained by each PC, then plot this for each PC.

```
pve <- pr.var/sum(pr.var)
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



To find specific data, we can call the rotation portion of our results, finding that the PC1 component for the variable “concave.points_mean” is -0.26. Going back to the summary of results, we can also see that a minimum of 5 PCs brings us to 80% variance.

```
wisc.pr$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21

Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

```
sum(pve[1:5])
```

```
[1] 0.8473427
```

Heirarchical Clustering


Now, let's perform some h-clustering on the data to see how that algorithm handles this data. First, we prepare the dataset by scaling it and finding distances. Then, we can find the `hclust()`.

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist,method="complete")
```

Plotting the results, we find that the height where there are 4 clusters is at 19.

```
plot(wisc.hclust)
abline(h=19,col="red",lty=2)
```

Cluster Dendrogram



```
data.dist
hclust (*, "complete")
```

Judging from the dendrogram, it may be worthwhile to cut the tree into 4 clusters. We can then use the `table()` function to compare against the diagnoses.

```
wisc.hclust.clusters <- cutree(wisc.hclust,k=4)
table(wisc.hclust.clusters,diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Repeating this for different cluster counts from 2 to 10 shows that 4 is probably the best cluster vs diagnosis match in terms of ratios. Below is an example of one of these repeats.

```
wisc.hclust.clusters8 <- cutree(wisc.hclust,k=8)
table(wisc.hclust.clusters8,diagnosis)
```

	diagnosis	
wisc.hclust.clusters8	B	M

1	12	86
2	0	79
3	0	3
4	331	39
5	2	0
6	12	1
7	0	2
8	0	2

We can also test this using the different h-clust methods available to see if any will fit better than the default "complete" method. Below is the code for the exploration of these alternative methods.

```
wisc.hclust.s <- hclust(data.dist,method="single")
wisc.hclust.s.clusters <- cutree(wisc.hclust.s,k=4)
table(wisc.hclust.s.clusters,diagnosis)
```

	diagnosis	
wisc.hclust.s.clusters	B	M
1	356	209
2	1	0
3	0	2
4	0	1

```
wisc.hclust.a <- hclust(data.dist,method="average")
wisc.hclust.a.clusters <- cutree(wisc.hclust.a,k=5)
table(wisc.hclust.a.clusters,diagnosis)
```

	diagnosis	
wisc.hclust.a.clusters	B	M
1	355	208
2	2	0
3	0	1
4	0	2
5	0	1

```
wisc.hclust.w <- hclust(data.dist,method="ward.D2")
wisc.hclust.w.clusters <- cutree(wisc.hclust.w,k=2)
table(wisc.hclust.w.clusters,diagnosis)
```


	diagnosis	
wisc.hclust.w.clusters	B	M
1	20	164
2	337	48

Of these new methods, ward.D2 seems to be the best, but still is not as representative of the diagnoses as the complete method, when judging by the ratios of “inaccurate” vs “accurate” data points.

Combining Methods

With both PCA and h-clustering, we can combine these methods to obtain a potentially better result of grouping. To do this, we can take the results from our PCA and perform h-clustering on it.

```
d <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(d,method="ward.D2")
plot(wisc.pr.hclust)
```




With 2 groups clearly delineated, we can cut the tree into 2.

```
grps <- cutree(wisc.pr.hclust,k=2)
head(grps)
```

```
842302    842517 84300903 84348301 84358402    843786
      1         1         1         1         1         1
```

Finally, we can plot our PC1 vs PC2, colored by the groups found in our previous step.

```
plot(wisc.pr$x[,1],wisc.pr$x[,2],col=grps)
```



To check against our expert diagnoses, we call the `table()` function.

```
table(grps,diagnosis)
```

```
      diagnosis
grps    B    M
1    24 179
2   333  33
```

Examining our results, we find that this combination of methods obtained the most “accurate” groups according to the ratios. Comparing to the other methods of heirarchical clustering and PCA, a combination of the two produced the best results.